

Latency — remote end-to-end vs. local compute are NOT comparable

